

# 一、centos7编译安装nginx 1.16.1稳定版

## 1.1 下载地址

已上传140SVN

```
http://39.100.254.140:12011/loit-Infrastructure-doc/loit-initproject-doc/blob/master/3%E3%80%81other/tools/nginx-1.16.1.tar.gz
```

```
http://39.100.254.140:12011/loit-Infrastructure-doc/loit-initproject-doc/blob/master/3%E3%80%81other/tools/echo-nginx-module-0.61.tar.gz
```

nginx-1.16.1.tar.gz 上传到目录: /usr/local/src echo-nginx-module-0.61.tar.gz 上传到root/echo-nginx-module-0.61.tar.gz 并解压

```
tar -zxvf echo-nginx-module-0.61.tar.gz
```

## 1.2 安装nginx

安装nginx 编译所需的lib库

```
yum -y install make zlib zlib-devel gcc-c++ libtool openssl openssl-devel  
yum -y install pcre pcre-devel
```

查看pcre(正则库)版本

```
pcre-config --version
```

```
8.32
```

进入编译目录

```
cd /usr/local/src
```

解压nginx压缩包

```
tar -zxvf nginx-1.16.1.tar.gz
```

进入解压目录

```
cd nginx-1.16.1
```

运行配置脚本(--prefix参数指定nginx安装的目录,默认安装在/usr/local/nginx)

```
./configure --prefix=/usr/local/nginx --add-module=/root/echo-nginx-module-0.61 --with-http_stub_status_module
```

编译安装nginx

```
make && make install
```

将nginx执行命令软链接到 /usr/bin

```
ln -s /usr/local/nginx/sbin/nginx /usr/bin
```

启动nginx

```
nginx
```

设置开机自启动

```
echo "/usr/local/nginx/sbin/nginx" >> /etc/rc.d/rc.local  
chmod +x /etc/rc.d/rc.local
```

### 1.3 测试echo模块

```
location /hello {  
    default_type 'text/plain';  
    return 200 'hello!';  
}  
  
location /hello_echo {  
    default_type 'text/plain';  
    echo "hello, echo!";  
}
```

```
curl http://127.0.0.1/hello
```

### 1.4 stub\_status模块用法

提供了查看 Nginx 运行的基本状态信息，我们只想让部分 IP 的人可以查看，此时可以配置一个访问控制：

```
vi /usr/local/nginx/conf/nginx.conf
```

```
location /nginx-status {
    stub_status;
    access_log off;
    #allow 192.168.179.0/24;
    #deny all;
}
```

```
curl 127.0.0.1/nginx-status
```

## 1.4 nginx相关命令

执行 `nginx -h` 查看相关命令

```
[root@localhost ~]# nginx -h
nginx version: nginx/1.16.1
Usage: nginx [-?hvVtTq] [-s signal] [-c filename] [-p prefix] [-g directives]

Options:
  -?, -h      : this help
  -v          : show version and exit
  -V          : show version and configure options then exit
  -t          : test configuration and exit
  -T          : test configuration, dump it and exit
  -q          : suppress non-error messages during configuration testing
  -s signal   : send signal to a master process: stop, quit, reopen, reload
  -p prefix   : set prefix path (default: /usr/local/nginx/)
  -c filename : set configuration file (default: conf/nginx.conf)
  -g directives : set global directives out of configuration file
```

复制代码

查看nginx安装目录

```
whereis nginx
```

停止重启

启动

```
[root sbin]# ./nginx
```

停止

```
[root sbin]# ./nginx -s stop
```

重启

```
[root sbin]# ./nginx -s reload
```

开启端口80

```
firewall-cmd --zone=public --add-port=80/tcp --permanent
```

命令含义：

-zone #作用域

-add-port=80/tcp #添加端口，格式为：端口/通讯协议

-permanent #永久生效，没有此参数重启后失效

重启防火墙

```
firewall-cmd --reload #重启firewall
systemctl stop firewalld.service #停止firewall
systemctl disable firewalld.service #禁止firewall开机启动
firewall-cmd --state #查看默认防火墙状态（关闭后显示notrunning，开启后显示running）
```

## 1.5 测试工具

### 1、ab 测试工具安装

```
yum -y install httpd-tools
```

测试2000连接数，50000次请求

```
ab -c 2000 -n 50000 http://172.16.203.78/hello
```

### 2、wrk测试工具

```
wrk -t50 -c300 -d30s -T30 http://172.16.203.78/hello
```

## 二、性能问题

### 2.1 未优化前

测试环境

```
虚拟机
cpu 核数：8核
内存：4G
```

查看当前cpu的状态：

```
[root ~]# lscpu |grep "CPU(s) "
```

## 🚗 1000 并发

```
wrk -t50 -c1000 -d30s -T30 http://172.16.203.78/hello

Running 30s test @ http://192.168.66.52/hello
50 threads and 1000 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
Latency       38.67ms  229.80ms   6.47s   97.16%
Req/Sec       1.90k    1.46k    17.12k   76.88%
2609441 requests in 30.08s, 380.63MB read
Socket errors: connect 29, read 0, write 0, timeout 0
Requests/sec: 86757.38
Transfer/sec: 12.65MB
```

- QPS 为 **86757.38**
- 平均延迟为 38.67ms
- 其中出现 **Socket errors**: connect 29

← → ↻ ⓘ 不安全 | 192.168.66.52/nginx-status

```
Active connections: 162
server accepts handled requests
307665 307665 13155971
Reading: 0 Writing: 1 Waiting: 161
```

- Active connections: 在200左右比较低
- Waiting数量比较多

```
top - 10:19:25 up 1 day, 6:01, 2 users, load average: 0.10, 0.13, 0.24
Tasks: 123 total, 2 running, 121 sleeping, 0 stopped, 0 zombie
%Cpu(s): 0.5 us, 25.6 sy, 0.0 ni, 73.6 id, 0.0 wa, 0.0 hi, 0.4 si, 0.0 st
KiB Mem : 3586780 total, 495896 free, 186280 used, 2904604 buff/cache
KiB Swap: 1572860 total, 1572860 free, 0 used. 3021912 avail Mem
```

| PID   | USER   | PR | NI | VIRT  | RES  | SHR  | S | %CPU  | %MEM | TIME+   | COMMAND      |
|-------|--------|----|----|-------|------|------|---|-------|------|---------|--------------|
| 44039 | nobody | 20 | 0  | 21188 | 1868 | 868  | R | 100.0 | 0.1  | 6:16.03 | nginx        |
| 9     | root   | 20 | 0  | 0     | 0    | 0    | S | 0.3   | 0.0  | 0:04.40 | rcu_sched    |
| 18    | root   | 20 | 0  | 0     | 0    | 0    | S | 0.3   | 0.0  | 0:08.03 | ksoftirqd/2  |
| 430   | root   | 20 | 0  | 0     | 0    | 0    | S | 0.3   | 0.0  | 0:02.27 | xfsaild/dm-0 |
| 1     | root   | 20 | 0  | 46064 | 6272 | 3892 | S | 0.0   | 0.2  | 0:07.12 | systemd      |
| 2     | root   | 20 | 0  | 0     | 0    | 0    | S | 0.0   | 0.0  | 0:00.05 | kthreadd     |

- CPU使用率 100%
- 内存0.1%

## 🚗 1000 并发

```
wrk -t50 -c2000 -d30s -T30 http://192.168.66.52/hello
unable to create thread 28: Too many open files
```

## 查看nginx错误日志

```
# tail /usr/local/nginx/logs/error.log
```

```
2020/08/07 08:33:30 [error] 44039#0: *59630 open() "/usr/local/nginx/html/favicon.ico"
failed (2: No such file or directory), client: 192.168.66.240, server: localhost, request:
"GET /favicon.ico HTTP/1.1", host: "192.168.66.52"
2020/08/07 08:48:19 [crit] 44039#0: accept4() failed (24: Too many open files)
```

- 出现错误 Too many open files

## 2.2 问题总结

- 并发1000出现socket异常、nginx Active 数量少、Waiting数量多。
- 并发2000出现 Too many open files 异常。

## 三、优化思路

- 1、系统和nginx是否可以建立多个socket连接
- 2、系统和nginx是否允许一次性打开多个文件

### 建立socket连接，从操作系统和nginx两个层面分析

#### (1) 从nginx

- 1、http连接快速关闭即配置nginx的 keep\_alive\_time:0。因为在HTTP 1.0中协议是 请求-》连接-》断开，即每次请求之后都需要再次握手，但是随着web应用的丰富出现很多css文件和其他资源文件，这就使得要求是否一次请求可以请求多个文件，这就是HTTP 1.1。
- 2、子进程允许打开的连接即配置nginx的 (worker\_connections)

#### (2) 从系统层面：

- (1)修改最大连接数 somaxconn(具体路径在 /proc/sys/net/core/somaxconn) (2)加快tcp连接的回收，即修改 (/proc/sys/net/ipv4/tcp\_tw\_recycle) (3)空闲的tcp是否允许回收利用，即修改 (/proc/sys/net/ipv4/tcp\_tw\_reuse) (4)是否对洪水抵御做相应的cookie操作，修改 (/proc/sys/net/ipv4/tcp\_syncookies)

### 打开文件方面

1.nginx: 子进程允许打开的文件数量：配置添加：worker\_rlimit\_nofile 2.系统：设置ulimit -n 设置一个较大的值

#### 一、最大打开文件数的限制

```
vi /etc/security/limits.conf
```

#### 最后添加

```
# End of file
root soft nofile 65535
root hard nofile 65535
* soft nofile 65535
* hard nofile 65535
```

## 二、用户进程限制

```
vi /etc/security/limits.d/20-nproc.conf
```

```
#加大普通用户限制 也可以改为unlimited
*          soft    nproc    40960
root       soft    nproc    unlimited
```

## 四、优化

内核参数：

```
vi /etc/sysctl.conf
```

```
net.ipv4.conf.default.rp_filter = 1

net.ipv4.ip_forward = 1
net.ipv4.conf.default.accept_source_route = 0
kernel.sysrq = 0
kernel.core_uses_pid = 1
kernel.msgmnb = 65536
kernel.msgmax = 65536
kernel.shmmax = 68719476736
kernel.shmall = 4294967296
net.ipv4.ip_local_port_range = 1024 65535
net.ipv4.tcp_max_syn_backlog = 65535

net.ipv4.tcp_max_tw_buckets = 262144

net.core.somaxconn = 65535
net.core.netdev_max_backlog = 200000
net.core.rmem_default = 67108864
net.core.wmem_default = 67108864
net.core.rmem_max = 67108864
net.core.wmem_max = 67108864
net.ipv4.tcp_rmem = 4096 87380 6291456
```

```
net.ipv4.tcp_wmem = 4096 65536 4194304
net.ipv4.tcp_mem = 3097431 4129911 6194862
net.ipv4.tcp_timestamps = 0

net.ipv4.tcp_syncookies = 1

net.ipv4.tcp_synack_retries = 1
net.ipv4.tcp_syn_retries = 1
net.ipv4.tcp_tw_reuse = 1
net.ipv4.tcp_tw_recycle = 0
net.ipv4.ip_forward = 1
net.ipv4.tcp_fin_timeout = 15
net.ipv4.tcp_keepalive_time = 120
vm.overcommit_memory = 1
fs.file-max = 1048576
```

## 运行 sysctl -p 后配置生效

```
sysctl -p
```

## 部分参数说明

`net.ipv4.tcp_syncookies = 1` 表示开启SYN Cookies。当出现SYN等待队列溢出时，启用cookies来处理，可防范少量SYN攻击，默认为0，表示关闭；

`net.ipv4.tcp_tw_reuse = 1` 表示开启重用。允许将TIME-WAIT sockets重新用于新的TCP连接，默认为0，表示关闭；

`net.ipv4.tcp_fin_timeout = 720` 表示如果套接字由本端要求关闭，这个参数决定了它保持在FIN-WAIT-2状态的时间。

## Nginx 参考配置文件如下：

```
user root;
worker_processes 10;
#daemon off;
#master_process off;

worker_cpu_affinity
#0000000000001
#0000000000010
000000000100
000000001000
000000010000
000000100000
000001000000
000010000000
```



```

0001000000000
0010000000000
0100000000000
1000000000000
;

#error_log    logs/error.log debug;
error_log     logs/error.log;

worker_rlimit_core 200m;
working_directory /tmp;

pid           logs/nginx.pid;
events {
    worker_connections 204800;
    use epoll;
    accept_mutex off;
    multi_accept on;
}

http {
    sendfile            on;
    tcp_nodelay         on;
    tcp_nopush          on;

    access_log off;

    server_tokens off;
    reset_timedout_connection on;

    keepalive_timeout 120;
    keepalive_requests 100000;

    client_max_body_size 20m;
    client_body_buffer_size 1024k;
    client_body_temp_path /tmp;

    upstream redis_cluster {
        testupstream_node $node_ip;
        server 0.0.0.0;
        keepalive 1024;
    }

    server {
        listen 80;
        server_name localhost backlog=204800;

        set $backserver "redis_cluster";
        set $node_ip "";

        location ~* "^/hdp/kvstore/" {
            testupstream_pass $backserver;
            testupstream_next_upstream error timeout invalid_response;

```

```

    }

    location /hello {
        default_type 'text/plain';
        return 200 'hello!';
    }

    location /hello_echo {
        default_type 'text/plain';
        echo "hello, echo!";
    }

    location /nginx-status {
        stub_status;
        access_log off;
        #allow 192.168.179.0/24;
        #deny all;
    }
}
}

```

## 五、优化后测试

### 服务器配置

cpu 8核  
内存 32G

### 300 并发

```

[root@iZbp12plbi27m4rkrqor65Z wrk-master]# wrk -t50 -c300 -d30s -T30 http://172.16.203.78/hello
Running 30s test @ http://172.16.203.78/hello
 50 threads and 300 connections
   Thread Stats   Avg    Stdev   Max   +/-  Stdev
   Latency   17.26ms   41.92ms  413.75ms   90.38%
   Req/Sec   581.47    248.00    3.15k    70.82%
858493 requests in 30.10s, 121.17MB read
Requests/sec: 28523.93
Transfer/sec:  4.03MB

```

- QPS 为 **28523.93**
- 平均延迟为 17.26ms

## 🖥️ 5000 并发

```
[root@iZbp12plbi27m4rkrqor65Z wrk-master]# wrk -t50 -c5000 -d30s -T30 http://172.16.203.78/hello
Running 30s test @ http://172.16.203.78/hello
 50 threads and 5000 connections
   Thread Stats   Avg     Stdev     Max   +/-  Stdev
   Latency    122.32ms   353.06ms   9.04s   94.40%
   Req/Sec    617.03    147.82    2.89k    73.40%
 921864 requests in 30.10s, 130.12MB read
Requests/sec: 30625.38
Transfer/sec: 4.32MB
[root@iZbp12plbi27m4rkrqor65Z wrk-master]#
```

- QPS 为 **30625.28**
- 平均延迟为 122.32ms

## 🖥️ 8000 并发

```
[root@iZbp12plbi27m4rkrqor65Z wrk-master]#
[root@iZbp12plbi27m4rkrqor65Z wrk-master]# wrk -t50 -c8000 -d30s -T30 http://172.16.203.78/hello
Running 30s test @ http://172.16.203.78/hello
 50 threads and 8000 connections
   Thread Stats   Avg     Stdev     Max   +/-  Stdev
   Latency    219.62ms   602.03ms  21.07s   93.19%
   Req/Sec    585.58    183.34    7.19k    74.72%
 876482 requests in 30.10s, 123.71MB read
Socket errors: connect 0, read 7, write 0, timeout 0
Requests/sec: 29120.44
Transfer/sec: 4.11MB
[root@iZbp12plbi27m4rkrqor65Z wrk-master]#
```

[https://blog.csdn.net/qq\\_27384769](https://blog.csdn.net/qq_27384769)

- QPS 为 **29120.44**
- 平均延迟为 219.62ms
- socket errors: **read 7**

## 🖥️ 10000 并发

```
[root@iZbp12plbi27m4rkrqor65Z wrk-master]# wrk -t50 -c10000 -d30s -T30 http://172.16.203.78/hello
Running 30s test @ http://172.16.203.78/hello
 50 threads and 10000 connections
   Thread Stats   Avg     Stdev     Max   +/-  Stdev
   Latency    286.33ms   782.05ms  22.95s   92.85%
   Req/Sec    573.29    172.92    6.17k    74.42%
 858507 requests in 30.09s, 121.17MB read
Socket errors: connect 0, read 19, write 0, timeout 0
Requests/sec: 28526.80
Transfer/sec: 4.03MB
[root@iZbp12plbi27m4rkrqor65Z wrk-master]#
```

[https://blog.csdn.net/qq\\_27384769](https://blog.csdn.net/qq_27384769)

- QPS 为 **28526**
- 平均延迟为 286.33ms
- socket errors: **read 19**

## 🖥️ 20000 并发

```
[root@iZbp12plbi27m4rkrqor65Z wrk-master]#
[root@iZbp12plbi27m4rkrqor65Z wrk-master]# wrk -t50 -c20000 -d30s -T30 http://172.16.203.78/hello
Running 30s test @ http://172.16.203.78/hello
 50 threads and 20000 connections
   Thread Stats   Avg     Stdev     Max   +/-  Stdev
   Latency    524.10ms    1.36s   27.86s   92.81%
   Req/Sec    613.50    203.68    9.37k    78.37%
 920085 requests in 30.11s, 129.86MB read
Socket errors: connect 7, read 483, write 0, timeout 0
Requests/sec: 30553.62
Transfer/sec: 4.31MB
[root@iZbp12plbi27m4rkrqor65Z wrk-master]#
```

[https://blog.csdn.net/qq\\_27384769](https://blog.csdn.net/qq_27384769)

- QPS 为 **30553.62**

- 平均延迟为 286.33ms
- socket errors: **connect 7 read 483**

#### CPU使用情况

```
%Cpu(s): 16.3 us, 20.1 sy, 0.0 ni, 52.1 id, 2.2 wa, 0.0 hi, 9.3 si, 0.0 st
KiB Mem : 32779824 total, 229200 free, 12678868 used, 19871756 buff/cache
KiB Swap: 0 total, 0 free, 0 used. 19643476 avail Mem
```

| PID   | USER   | PR | NI  | VIRT    | RES    | SHR   | S | %CPU | %MEM | TIME+     | COMMAND      |
|-------|--------|----|-----|---------|--------|-------|---|------|------|-----------|--------------|
| 6843  | root   | 20 | 0   | 6325632 | 1.3g   | 7640  | S | 59.8 | 4.1  | 145:29.98 | java         |
| 3     | root   | 20 | 0   | 0       | 0      | 0     | S | 39.2 | 0.0  | 4:17.31   | ksoftirqd/0  |
| 8151  | root   | 20 | 0   | 8057404 | 2.5g   | 7480  | S | 33.9 | 7.9  | 1437:23   | java         |
| 28298 | esroot | 20 | 0   | 12.8g   | 1.5g   | 11276 | S | 32.6 | 4.9  | 126:41.09 | java         |
| 994   | nobody | 20 | 0   | 45568   | 26224  | 828   | S | 31.6 | 0.1  | 1:17.38   | nginx        |
| 996   | nobody | 20 | 0   | 45568   | 26224  | 828   | S | 30.9 | 0.1  | 1:19.36   | nginx        |
| 995   | nobody | 20 | 0   | 45592   | 26376  | 864   | S | 29.9 | 0.1  | 1:13.97   | nginx        |
| 992   | nobody | 20 | 0   | 45592   | 26380  | 864   | S | 29.6 | 0.1  | 1:08.67   | nginx        |
| 997   | nobody | 20 | 0   | 45568   | 26224  | 828   | S | 28.9 | 0.1  | 1:11.14   | nginx        |
| 991   | nobody | 20 | 0   | 45568   | 26196  | 800   | S | 27.6 | 0.1  | 1:04.99   | nginx        |
| 990   | nobody | 20 | 0   | 45592   | 26380  | 864   | S | 27.2 | 0.1  | 1:06.47   | nginx        |
| 993   | nobody | 20 | 0   | 45568   | 26224  | 828   | S | 26.6 | 0.1  | 1:03.88   | nginx        |
| 2257  | root   | 10 | -10 | 147048  | 27044  | 3448  | S | 2.0  | 0.1  | 379:52.07 | AliYunDun    |
| 817   | root   | 0  | -20 | 0       | 0      | 0     | S | 1.3  | 0.0  | 0:52.87   | kworker/0:1H |
| 8494  | root   | 20 | 0   | 6388220 | 546028 | 7504  | S | 1.3  | 1.7  | 158:08.72 | java         |

- 测试过程平均cpu使用率40%

## 六、nginx性能测试结论

nginx优化的方法三种，第一种优化linux内核参数，使内核变的更为强大，第二种是优化nginx配置文件，使nginx变的更为强大，第三种是扩展服务器的cpu和内存，使服务器变的更为强大。

#### 单机测试：

- 单机8核cpu的平均在30000QPS, 1万并发连接数平均消耗40%cpu。
- nginx并发数与**cpu核数**有关，cpu核数到达**88核**可以实现百万QPS数量。
- 并发连接数达到8000 ~ 10000 开始有很少量的error，并发连接数达到20000 error 数量开始上升。

#### 参考内存配置要求：

在操作系统层面每个TCP连接会占用3k-10k的内存，以20万来计算，需要2G内存。nginx程序本身还要消耗内存，特别是nginx反向代理POST请求比较多的情况，20万连接情况下推荐16G内存配置。